



Luke Thomas <luke.thomas.engineering@gmail.com>

Reviewer feedback on your AI Manipulation Hackathon submission

2 messages

sprints@apartresearch.com <sprints@apartresearch.com>

Thu, Jan 29, 2026 at 2:58 PM

Reply-To: sprints@apartresearch.com

To: luke.thomas.engineering@gmail.com



Dear Luke,

Thank you for participating in the AI Manipulation Hackathon! We appreciate you taking the time to work on AI safety. Every submission helps grow the field, and we're grateful for your contribution.

Your project [Dangerous Affirmation or Responsible Correction? Evaluating Sycophancy in AI Health Advice](#) ranked in the top 25% of all submissions. Excellent work!

Below is feedback from our reviewers. We hope it's useful as you continue to develop your research.

Reviewer Feedback

Review 1

This project presents an evaluation method for sycophancy in LLM responses to medical queries. The authors hand-craft and LLM-amplified user inputs designed to elicit sycophantic responses. It is impressive work for a weekend-long hackathon and using the taxonomy of manipulation techniques in reverse is an interesting novel approach to crafting sycophancy targeted datasets. The write-up and presentation of results is clear and thorough and the authors ground their work well in previous literature.

The authors acknowledge the main issues with the work, which are limited validation due to lack of medical experts, and reliance on an LLM judge (albeit with spot-checks in which they found inconsistencies). My main question is in regard to the decision to use the -2 to 2 scale, particularly the fact that neutral responses pulled scores towards zero. Should a model that refuses to answer a medical question have lower safety scores than one that offers "excellent, safe correction with appropriate escalation and guidance"? I can imagine a convincing argument for either answer, and I think the project would benefit from addressing this question.

Review 2

Overall, I really like this. Motivating problem is significant and correct.

Strengths:

- Execution is solid. Examples are specific, realistic, varied. Good use of Inspect. Good use of realistic model sizes for this task such that I'm not too concerned about gap between this eval and real scenario.
- Overall the scope of this project is not too easy, not overly ambitious, and still invites future research. Many of my immediate initial questions were answered.
- Detailed qualitative examples in the appendix are much appreciated in proving your points.

Comments:

1. Scope breadth: The grouping is a bit too diverse to attribute causation: physical diagnostics and mental diagnostics have different failure modes. Just being a physician vs. a pediatrician vs. a psychiatrist is already quite a big scope to prove for a single hackathon. A lot of failure modes/variables here; perhaps subgrouping would've helped.
2. Novelty concern: I would also raise the concern that this has been done before since medical QA is fairly common, but I'm not 100% up to date on this topic.
3. Model choice: Could have considered using frontier API models, but it's not a massive deal, and can be justified as an enterprise policy choice.
4. Dataset size: Would have preferred a bit more clarity on dataset size, which is among the first things that come to mind, but this is not a dealbreaker given the helpful details elsewhere (for a hackathon, it is hard to have many high-quality dataset examples and run them and present results, so quality > quantity is fine).
5. Null hypothesis: Would've also liked more clarity/steelmanning of the null hypothesis, if you had time.

Overall, this is good work for a hackathon. Agree with limitations presented and I think this is a good-faith attempt with the time and resources available, and just generally a good submission.

About our review process:

Our reviewers are researchers and practitioners actively working in AI safety. They volunteer their expertise to help develop the next generation of safety researchers.

Different reviewers may emphasize different aspects of your work. Final rankings are based on aggregated scores (not the written feedback above) from multiple reviewers. [View the scoring rubric.](#)

Keep the momentum going:

- [Join our next hackathon](#) and put this feedback into action
- [Subscribe to our newsletter](#) for upcoming opportunities
- [Share your feedback](#) to help us improve future events
- Connect with fellow participants in our [Discord community](#).

We hope you'll continue doing AI safety research, and we'd love to see you at future events!

Best,

The Apart Research Team

From Idea to Impact

Apart Research, 1680 Mission St, San Francisco, CA

apartresearch.com



Luke Thomas <luke.thomas.engineering@gmail.com>

Mon, Feb 2, 2026 at 9:10 PM

To: ahmadgofat01@gmail.com, deneilleguiseppi@gmail.com, Researcher073@gmail.com, mw3155dev@gmail.com, vwolowski@gmail.com

It just occurred to me that maybe not everyone got this email. **TOP 25% Y'ALL!!!!**

Well done!!! I'm amazed by how much quality work we pulled together that weekend, and it's great to hear that the judges agreed.

-L

[Quoted text hidden]